

Go to the [first](#), [previous](#), [next](#), [last](#) section, [table of contents](#).

Vectors and Matrices

The functions described in this chapter provide a simple vector and matrix interface to ordinary C arrays. The memory management of these arrays is implemented using a single underlying type, known as a block. By writing your functions in terms of vectors and matrices you can pass a single structure containing both data and dimensions as an argument without needing additional function parameters. The structures are compatible with the vector and matrix formats used by BLAS routines.

Data types

All the functions are available for each of the standard data-types. The versions for `double` have the prefix `gsl_block`, `gsl_vector` and `gsl_matrix`. Similarly the versions for single-precision `float` arrays have the prefix `gsl_block_float`, `gsl_vector_float` and `gsl_matrix_float`. The full list of available types is given below,

<code>gsl_block</code>	<code>double</code>
<code>gsl_block_float</code>	<code>float</code>
<code>gsl_block_long_double</code>	<code>long double</code>
<code>gsl_block_int</code>	<code>int</code>
<code>gsl_block_uint</code>	<code>unsigned int</code>
<code>gsl_block_long</code>	<code>long</code>
<code>gsl_block_ulong</code>	<code>unsigned long</code>
<code>gsl_block_short</code>	<code>short</code>
<code>gsl_block_ushort</code>	<code>unsigned short</code>
<code>gsl_block_char</code>	<code>char</code>
<code>gsl_block_uchar</code>	<code>unsigned char</code>
<code>gsl_block_complex</code>	<code>complex double</code>
<code>gsl_block_complex_float</code>	<code>complex float</code>
<code>gsl_block_complex_long_double</code>	<code>complex long double</code>

Corresponding types exist for the `gsl_vector` and `gsl_matrix` functions.

Blocks

For consistency all memory is allocated through a `gsl_block` structure. The structure contains two components, the size of an area of memory and a pointer to the memory. The `gsl_block` structure looks like this,

```
typedef struct
{
    size_t size;
    double * data;
} gsl_block;
```

Vectors and matrices are made by **slicing** an underlying block. A slice is a set of elements formed from an initial offset and a combination of indices and step-sizes. In the case of a matrix the step-size for the column index represents the row-length. The step-size for a vector is known as the **stride**.

The functions for allocating and deallocating blocks are defined in ``gsl_block.h'`

Block allocation

The functions for allocating memory to a block follow the style of `malloc` and `free`. In addition they also perform their own error checking. If there is insufficient memory available to allocate a block then the functions call the GSL error handler (with an error number of `GSL_ENOMEM`) in addition to returning a null pointer. Thus if you use the library error handler to abort your program then it isn't necessary to check every `alloc`.

Function: `gsl_block * gsl_block_alloc (size_t n)`

This function allocates memory for a block of n double-precision elements, returning a pointer to the block struct. The block is not initialized and so the values of its elements are undefined. Use the function `gsl_block_calloc` if you want to ensure that all the elements are initialized to zero.

A null pointer is returned if insufficient memory is available to create the block.

Function: `gsl_block * gsl_block_calloc (size_t n)`

This function allocates memory for a block and initializes all the elements of the block to zero.

Function: `void gsl_block_free (gsl_block * b)`

This function frees the memory used by a block b previously allocated with `gsl_block_alloc` or `gsl_block_calloc`.

Reading and writing blocks

The library provides functions for reading and writing blocks to a file as binary data or formatted text.

Function: `int gsl_block_fwrite (FILE * stream, const gsl_block * b)`

This function writes the elements of the block b to the stream $stream$ in binary format. The return value is 0 for success and `GSL_EFAILED` if there was a problem writing to the file. Since the data is written in the native binary format it may not be portable between different architectures.

Function: `int gsl_block_fread (FILE * stream, gsl_block * b)`

This function reads into the block b from the open stream $stream$ in binary format. The block b must be preallocated with the correct length since the function uses the size of b to determine how many bytes to read. The return value is 0 for success and `GSL_EFAILED` if there was a problem reading from the file. The data is assumed to have been written in the native binary format on the same architecture.

Function: `int gsl_block_fprintf (FILE * stream, const gsl_block * b, const char * format)`

This function writes the elements of the block b line-by-line to the stream $stream$ using the format specifier $format$, which should be one of the `%g`, `%e` or `%f` formats for floating point numbers and `%d` for integers. The function returns 0 for success and `GSL_EFAILED` if there was a problem writing to the file.

Function: `int gsl_block_fscanf (FILE * stream, gsl_block * b)`

This function reads formatted data from the stream *stream* into the block *b*. The block *b* must be preallocated with the correct length since the function uses the size of *b* to determine how many numbers to read. The function returns 0 for success and `GSL_EFAILED` if there was a problem reading from the file.

Example programs for blocks

The following program shows how to allocate a block,

```
#include <stdio.h>
#include <gsl/gsl_block.h>

int
main (void)
{
    gsl_block * b = gsl_block_alloc (100);

    printf("length of block = %u\n", b->size);
    printf("block data address = %#x\n", b->data);

    gsl_block_free (b);
    return 0;
}
```

Here is the output from the program,

```
length of block = 100
block data address = 0x804b0d8
```

Vectors

Vectors are defined by a `gsl_vector` structure which describes a slice of a block. Different vectors can be created which point to the same block. A vector slice is a set of equally-spaced elements of an area of memory.

The `gsl_vector` structure contains five components, the **size**, the **stride**, a pointer to the memory where the elements are stored, *data*, a pointer to the block owned by the vector, *block*, if any, and an ownership flag, *owner*. The structure is very simple and looks like this,

```
typedef struct
{
    size_t size;
    size_t stride;
    double * data;
    gsl_block * block;
    int owner;
} gsl_vector;
```

The *size* is simply the number of vector elements. The range of valid indices runs from 0 to *size*-1.

The *stride* is the step-size from one element to the next in physical memory, measured in units of the appropriate datatype. The pointer *data* gives the location of the first element of the vector in memory. The pointer *block* stores the location of the memory block in which the vector elements are located (if any). If the vector owns this block then the *owner* field is set to one and the block will be deallocated when the vector is freed. If the vector points to a block owned by another object then the *owner* field is zero and any underlying block will not be deallocated.

The functions for allocating and accessing vectors are defined in ``gsl_vector.h'`

Vector allocation

The functions for allocating memory to a vector follow the style of `malloc` and `free`. In addition they also perform their own error checking. If there is insufficient memory available to allocate a vector then the functions call the GSL error handler (with an error number of `GSL_ENOMEM`) in addition to returning a null pointer. Thus if you use the library error handler to abort your program then it isn't necessary to check every `alloc`.

Function: `gsl_vector * gsl_vector_alloc (size_t n)`

This function creates a vector of length *n*, returning a pointer to a newly initialized vector struct. A new block is allocated for the elements of the vector, and stored in the *block* component of the vector struct. The block is "owned" by the vector, and will be deallocated when the vector is deallocated.

Function: `gsl_vector * gsl_vector_calloc (size_t n)`

This function allocates memory for a vector of length *n* and initializes all the elements of the vector to zero.

Function: `void gsl_vector_free (gsl_vector * v)`

This function frees a previously allocated vector *v*. If the vector was created using `gsl_vector_alloc` then the block underlying the vector will also be deallocated. If the vector has been created from another object then the memory is still owned by that object and will not be deallocated.

Accessing vector elements

Unlike FORTRAN compilers, C compilers do not usually provide support for range checking of vectors and matrices. Range checking is available in the GNU C Compiler extension `checkgcc` but it is not available on every platform. The functions `gsl_vector_get` and `gsl_vector_set` can perform portable range checking for you and report an error if you attempt to access elements outside the allowed range.

The functions for accessing the elements of a vector or matrix are defined in ``gsl_vector.h'` and declared `extern inline` to eliminate function-call overhead. If necessary you can turn off range checking completely without modifying any source files by recompiling your program with the preprocessor definition `GSL_RANGE_CHECK_OFF`. Provided your compiler supports inline functions the effect of turning off range checking is to replace calls to `gsl_vector_get(v, i)` by `v->data[i*v->stride]` and calls to `gsl_vector_set(v, i, x)` by `v->data[i*v->stride]=x`. Thus there should be no performance penalty for using the range checking functions when range checking is turned off.

Function: double **gsl_vector_get** (*const gsl_vector * v, size_t i*)

This function returns the i -th element of a vector v . If i lies outside the allowed range of 0 to $n-1$ then the error handler is invoked and 0 is returned.

Function: void **gsl_vector_set** (*gsl_vector * v, size_t i, double x*)

This function sets the value of the i -th element of a vector v to x . If i lies outside the allowed range of 0 to $n-1$ then the error handler is invoked.

Function: double * **gsl_vector_ptr** (*gsl_vector * v, size_t i*)

Function: const double * **gsl_vector_ptr** (*const gsl_vector * v, size_t i*)

These functions return a pointer to the i -th element of a vector v . If i lies outside the allowed range of 0 to $n-1$ then the error handler is invoked and a null pointer is returned.

Initializing vector elements

Function: void **gsl_vector_set_all** (*gsl_vector * v, double x*)

This function sets all the elements of the vector v to the value x .

Function: void **gsl_vector_set_zero** (*gsl_vector * v*)

This function sets all the elements of the vector v to zero.

Function: int **gsl_vector_set_basis** (*gsl_vector * v, size_t i*)

This function makes a basis vector by setting all the elements of the vector v to zero except for the i -th element which is set to one.

Reading and writing vectors

The library provides functions for reading and writing vectors to a file as binary data or formatted text.

Function: int **gsl_vector_fwrite** (*FILE * stream, const gsl_vector * v*)

This function writes the elements of the vector v to the stream $stream$ in binary format. The return value is 0 for success and `GSL_EFAILED` if there was a problem writing to the file. Since the data is written in the native binary format it may not be portable between different architectures.

Function: int **gsl_vector_fread** (*FILE * stream, gsl_vector * v*)

This function reads into the vector v from the open stream $stream$ in binary format. The vector v must be preallocated with the correct length since the function uses the size of v to determine how many bytes to read. The return value is 0 for success and `GSL_EFAILED` if there was a problem reading from the file. The data is assumed to have been written in the native binary format on the same architecture.

Function: int **gsl_vector_fprintf** (*FILE * stream, const gsl_vector * v, const char * format*)

This function writes the elements of the vector v line-by-line to the stream $stream$ using the format specifier $format$, which should be one of the `%g`, `%e` or `%f` formats for floating point numbers and `%d` for integers. The function returns 0 for success and `GSL_EFAILED` if there was a problem writing to the file.

Function: int **gsl_vector_fscanf** (*FILE * stream, gsl_vector * v*)

This function reads formatted data from the stream *stream* into the vector *v*. The vector *v* must be preallocated with the correct length since the function uses the size of *v* to determine how many numbers to read. The function returns 0 for success and `GSL_EFAILED` if there was a problem reading from the file.

Vector views

In addition to creating vectors from slices of blocks it is also possible to slice vectors and create vector views. For example, a subvector of another vector can be described with a view, or two views can be made which provide access to the even and odd elements of a vector.

A vector view is a temporary object, stored on the stack, which can be used to operate on a subset of vector elements. Vector views can be defined for both constant and non-constant vectors, using separate types that preserve constness. A vector view has the type `gsl_vector_view` and a constant vector view has the type `gsl_vector_const_view`. In both cases the elements of the view can be accessed as a `gsl_vector` using the `vector` component of the view object. A pointer to a vector of type `gsl_vector *` or `const gsl_vector *` can be obtained by taking the address of this component with the `&` operator.

Function: `gsl_vector_view` **`gsl_vector_subvector`** (*gsl_vector *v, size_t offset, size_t n*)

Function: `gsl_vector_const_view` **`gsl_vector_const_subvector`** (*const gsl_vector *v, size_t offset, size_t n*)

These functions return a vector view of a subvector of another vector *v*. The start of the new vector is offset by *offset* elements from the start of the original vector. The new vector has *n* elements.

Mathematically, the *i*-th element of the new vector *v'* is given by,

$$v'(i) = v->data[(offset + i)*v->stride]$$

where the index *i* runs from 0 to *n*-1.

The `data` pointer of the returned vector struct is set to null if the combined parameters (*offset,n*) overrun the end of the original vector.

The new vector is only a view of the block underlying the original vector, *v*. The block containing the elements of *v* is not owned by the new vector. When the view goes out of scope the original vector *v* and its block will continue to exist. The original memory can only be deallocated by freeing the original vector. Of course, the original vector should not be deallocated while the view is still in use.

The function `gsl_vector_const_subvector` is equivalent to `gsl_vector_subvector` but can be used for vectors which are declared `const`.

Function: `gsl_vector` **`gsl_vector_subvector_with_stride`** (*gsl_vector *v, size_t offset, size_t stride, size_t n*)

Function: `gsl_vector_const_view` **`gsl_vector_const_subvector_with_stride`** (*const gsl_vector *v, size_t offset, size_t stride, size_t n*)

These functions return a vector view of a subvector of another vector *v* with an additional stride argument. The subvector is formed in the same way as for `gsl_vector_subvector` but the new vector has *n* elements with a step-size of *stride* from one element to the next in the original vector.

Mathematically, the *i*-th element of the new vector *v'* is given by,

$$v'(i) = v->data[(offset + i*stride)*v->stride]$$

where the index i runs from 0 to $n-1$.

Note that subvector views give direct access to the underlying elements of the original vector. For example, the following code will zero the even elements of the vector v of length n , while leaving the odd elements untouched,

```
gsl_vector_view v_even
  = gsl_vector_subvector_with_stride (v, 0, 2, n/2);
gsl_vector_set_zero (&v_even.vector);
```

A vector view can be passed to any subroutine which takes a vector argument just as a directly allocated vector would be, using `&view.vector`. For example, the following code computes the norm of odd elements of v using the BLAS routine DNRM2,

```
gsl_vector_view v_odd
  = gsl_vector_subvector_with_stride (v, 1, 2, n/2);
double r = gsl_blas_dnorm2 (&v_odd.vector);
```

The function `gsl_vector_const_subvector_with_stride` is equivalent to `gsl_vector_subvector_with_stride` but can be used for vectors which are declared `const`.

Function: `gsl_vector_view` **`gsl_vector_complex_real`** (*`gsl_vector_complex *v`*)

Function: `gsl_vector_const_view` **`gsl_vector_complex_const_real`** (*`const gsl_vector_complex *v`*)

These functions return a vector view of the real parts of the complex vector v .

The function `gsl_vector_complex_const_real` is equivalent to `gsl_vector_complex_real` but can be used for vectors which are declared `const`.

Function: `gsl_vector_view` **`gsl_vector_complex_imag`** (*`gsl_vector_complex *v`*)

Function: `gsl_vector_const_view` **`gsl_vector_complex_const_imag`** (*`const gsl_vector_complex *v`*)

These functions return a vector view of the imaginary parts of the complex vector v .

The function `gsl_vector_complex_const_imag` is equivalent to `gsl_vector_complex_imag` but can be used for vectors which are declared `const`.

Function: `gsl_vector_view` **`gsl_vector_view_array`** (*`double *base, size_t n`*)

Function: `gsl_vector_const_view` **`gsl_vector_const_view_array`** (*`const double *base, size_t n`*)

These functions return a vector view of an array. The start of the new vector is given by $base$ and has n elements. Mathematically, the i -th element of the new vector v' is given by,

$$v'(i) = base[i]$$

where the index i runs from 0 to $n-1$.

The array containing the elements of v is not owned by the new vector view. When the view goes out of scope the original array will continue to exist. The original memory can only be deallocated by freeing the original pointer $base$. Of course, the original array should not be deallocated while the view is still in use.

The function `gsl_vector_const_view_array` is equivalent to `gsl_vector_view_array` but can be used for arrays which are declared `const`.

Function: `gsl_vector_view` **`gsl_vector_view_array_with_stride`** (*double * base, size_t stride, size_t n*)

Function: `gsl_vector_const_view` **`gsl_vector_const_view_array_with_stride`** (*const double * base, size_t stride, size_t n*)

These functions return a vector view of an array *base* with an additional stride argument. The subvector is formed in the same way as for `gsl_vector_view_array` but the new vector has *n* elements with a step-size of *stride* from one element to the next in the original array. Mathematically, the *i*-th element of the new vector *v'* is given by,

$$v'(i) = \text{base}[i * \text{stride}]$$

where the index *i* runs from 0 to *n*-1.

Note that the view gives direct access to the underlying elements of the original array. A vector view can be passed to any subroutine which takes a vector argument just as a directly allocated vector would be, using `&view.vector`.

The function `gsl_vector_const_view_array_with_stride` is equivalent to `gsl_vector_view_array_with_stride` but can be used for arrays which are declared `const`.

Copying vectors

Common operations on vectors such as addition and multiplication are available in the BLAS part of the library (see section [BLAS Support](#)). However, it is useful to have a small number of utility functions which do not require the full BLAS code. The following functions fall into this category.

Function: `int` **`gsl_vector_memcpy`** (*gsl_vector * dest, const gsl_vector * src*)

This function copies the elements of the vector *src* into the vector *dest*. The two vectors must have the same length.

Function: `int` **`gsl_vector_swap`** (*gsl_vector * v, gsl_vector * w*)

This function exchanges the elements of the vectors *v* and *w* by copying. The two vectors must have the same length.

Exchanging elements

The following function can be used to exchange, or permute, the elements of a vector.

Function: `int` **`gsl_vector_swap_elements`** (*gsl_vector * v, size_t i, size_t j*)

This function exchanges the *i*-th and *j*-th elements of the vector *v* in-place.

Function: `int` **`gsl_vector_reverse`** (*gsl_vector * v*)

This function reverses the order of the elements of the vector *v*.

Vector operations

The following operations are only defined for real vectors.

Function: `int gsl_vector_add (gsl_vector * a, const gsl_vector * b)`

This function adds the elements of vector b to the elements of vector a , $a'_i = a_i + b_i$. The two vectors must have the same length.

Function: `int gsl_vector_sub (gsl_vector * a, const gsl_vector * b)`

This function subtracts the elements of vector b from the elements of vector a , $a'_i = a_i - b_i$. The two vectors must have the same length.

Function: `int gsl_vector_mul (gsl_vector * a, const gsl_vector * b)`

This function multiplies the elements of vector a by the elements of vector b , $a'_i = a_i * b_i$. The two vectors must have the same length.

Function: `int gsl_vector_div (gsl_vector * a, const gsl_vector * b)`

This function divides the elements of vector a by the elements of vector b , $a'_i = a_i / b_i$. The two vectors must have the same length.

Function: `int gsl_vector_scale (gsl_vector * a, const double x)`

This function multiplies the elements of vector a by the constant factor x , $a'_i = x a_i$.

Function: `int gsl_vector_add_constant (gsl_vector * a, const double x)`

This function adds the constant value x to the elements of the vector a , $a'_i = a_i + x$.

Finding maximum and minimum elements of vectors

Function: `double gsl_vector_max (const gsl_vector * v)`

This function returns the maximum value in the vector v .

Function: `double gsl_vector_min (const gsl_vector * v)`

This function returns the minimum value in the vector v .

Function: `void gsl_vector_minmax (const gsl_vector * v, double * min_out, double * max_out)`

This function returns the minimum and maximum values in the vector v , storing them in min_out and max_out .

Function: `size_t gsl_vector_max_index (const gsl_vector * v)`

This function returns the index of the maximum value in the vector v . When there are several equal maximum elements then the lowest index is returned.

Function: `size_t gsl_vector_min_index (const gsl_vector * v)`

This function returns the index of the minimum value in the vector v . When there are several equal minimum elements then the lowest index is returned.

Function: `void gsl_vector_minmax_index (const gsl_vector * v, size_t * imin, size_t * imax)`

This function returns the indices of the minimum and maximum values in the vector v , storing them in $imin$ and $imax$. When there are several equal minimum or maximum elements then the lowest indices are returned.

Vector properties

Function: `int gsl_vector_isnull (const gsl_vector * v)`

This function returns 1 if all the elements of the vector *v* are zero, and 0 otherwise.

Example programs for vectors

This program shows how to allocate, initialize and read from a vector using the functions `gsl_vector_alloc`, `gsl_vector_set` and `gsl_vector_get`.

```
#include <stdio.h>
#include <gsl/gsl_vector.h>

int
main (void)
{
    int i;
    gsl_vector * v = gsl_vector_alloc (3);

    for (i = 0; i < 3; i++)
    {
        gsl_vector_set (v, i, 1.23 + i);
    }

    for (i = 0; i < 100; i++)
    {
        printf("v_%d = %g\n", i, gsl_vector_get (v, i));
    }

    return 0;
}
```

Here is the output from the program. The final loop attempts to read outside the range of the vector *v*, and the error is trapped by the range-checking code in `gsl_vector_get`.

```
v_0 = 1.23
v_1 = 2.23
v_2 = 3.23
gsl: vector_source.c:12: ERROR: index out of range
IOT trap/Abort (core dumped)
```

The next program shows how to write a vector to a file.

```
#include <stdio.h>
#include <gsl/gsl_vector.h>

int
main (void)
{
    int i;
    gsl_vector * v = gsl_vector_alloc (100);

    for (i = 0; i < 100; i++)
    {
```

```

    gsl_vector_set (v, i, 1.23 + i);
}

{
    FILE * f = fopen("test.dat", "w");
    gsl_vector_fprintf (f, v, "%.5g");
    fclose (f);
}
return 0;
}

```

After running this program the file `test.dat` should contain the elements of `v`, written using the format specifier `%.5g`. The vector could then be read back in using the function `gsl_vector_fscanf (f, v)` as follows:

```

#include <stdio.h>
#include <gsl/gsl_vector.h>

int
main (void)
{
    int i;
    gsl_vector * v = gsl_vector_alloc (10);

    {
        FILE * f = fopen("test.dat", "r");
        gsl_vector_fscanf (f, v);
        fclose (f);
    }

    for (i = 0; i < 10; i++)
    {
        printf("%g\n", gsl_vector_get(v, i));
    }

    return 0;
}

```

Matrices

Matrices are defined by a `gsl_matrix` structure which describes a generalized slice of a block. Like a vector it represents a set of elements in an area of memory, but uses two indices instead of one.

The `gsl_matrix` structure contains six components, the two dimensions of the matrix, a physical dimension, a pointer to the memory where the elements of the matrix are stored, *data*, a pointer to the block owned by the matrix *block*, if any, and an ownership flag, *owner*. The physical dimension determines the memory layout and can differ from the matrix dimension to allow the use of submatrices. The `gsl_matrix` structure is very simple and looks like this,

```

typedef struct
{
    size_t size1;
    size_t size2;
}

```

```

size_t tda;
double * data;
gsl_block * block;
int owner;
} gsl_matrix;

```

Matrices are stored in row-major order, meaning that each row of elements forms a contiguous block in memory. This is the standard "C-language ordering" of two-dimensional arrays. Note that FORTRAN stores arrays in column-major order. The number of rows is *size1*. The range of valid row indices runs from 0 to *size1*-1. Similarly *size2* is the number of columns. The range of valid column indices runs from 0 to *size2*-1. The physical row dimension *tda*, or **trailing dimension**, specifies the size of a row of the matrix as laid out in memory.

For example, in the following matrix *size1* is 3, *size2* is 4, and *tda* is 8. The physical memory layout of the matrix begins in the top left hand-corner and proceeds from left to right along each row in turn.

```

00 01 02 03 xx xx xx xx
10 11 12 13 xx xx xx xx
20 21 22 23 xx xx xx xx

```

Each unused memory location is represented by "xx". The pointer *data* gives the location of the first element of the matrix in memory. The pointer *block* stores the location of the memory block in which the elements of the matrix are located (if any). If the matrix owns this block then the *owner* field is set to one and the block will be deallocated when the matrix is freed. If the matrix is only a slice of a block owned by another object then the *owner* field is zero and any underlying block will not be freed.

The functions for allocating and accessing matrices are defined in ``gsl_matrix.h'`

Matrix allocation

The functions for allocating memory to a matrix follow the style of `malloc` and `free`. They also perform their own error checking. If there is insufficient memory available to allocate a vector then the functions call the GSL error handler (with an error number of `GSL_ENOMEM`) in addition to returning a null pointer. Thus if you use the library error handler to abort your program then it isn't necessary to check every `alloc`.

Function: `gsl_matrix * gsl_matrix_alloc (size_t n1, size_t n2)`

This function creates a matrix of size *n1* rows by *n2* columns, returning a pointer to a newly initialized matrix struct. A new block is allocated for the elements of the matrix, and stored in the *block* component of the matrix struct. The block is "owned" by the matrix, and will be deallocated when the matrix is deallocated.

Function: `gsl_matrix * gsl_matrix_calloc (size_t n1, size_t n2)`

This function allocates memory for a matrix of size *n1* rows by *n2* columns and initializes all the elements of the matrix to zero.

Function: `void gsl_matrix_free (gsl_matrix * m)`

This function frees a previously allocated matrix *m*. If the matrix was created using `gsl_matrix_alloc` then the block underlying the matrix will also be deallocated. If the matrix has been created from another object then the memory is still owned by that object and will not be deallocated.

Accessing matrix elements

The functions for accessing the elements of a matrix use the same range checking system as vectors. You turn off range checking by recompiling your program with the preprocessor definition `GSL_RANGE_CHECK_OFF`.

The elements of the matrix are stored in "C-order", where the second index moves continuously through memory. More precisely, the element accessed by the function `gsl_matrix_get(m, i, j)` and `gsl_matrix_set(m, i, j, x)` is

```
m->data[i * m->tda + j]
```

where *tda* is the physical row-length of the matrix.

Function: double **gsl_matrix_get** (*const gsl_matrix * m, size_t i, size_t j*)

This function returns the (i,j) th element of a matrix *m*. If *i* or *j* lie outside the allowed range of 0 to *n1*-1 and 0 to *n2*-1 then the error handler is invoked and 0 is returned.

Function: void **gsl_matrix_set** (*gsl_matrix * m, size_t i, size_t j, double x*)

This function sets the value of the (i,j) th element of a matrix *m* to *x*. If *i* or *j* lies outside the allowed range of 0 to *n1*-1 and 0 to *n2*-1 then the error handler is invoked.

Function: double * **gsl_matrix_ptr** (*gsl_matrix * m, size_t i, size_t j*)

Function: const double * **gsl_matrix_ptr** (*const gsl_matrix * m, size_t i, size_t j*)

These functions return a pointer to the (i,j) th element of a matrix *m*. If *i* or *j* lie outside the allowed range of 0 to *n1*-1 and 0 to *n2*-1 then the error handler is invoked and a null pointer is returned.

Initializing matrix elements

Function: void **gsl_matrix_set_all** (*gsl_matrix * m, double x*)

This function sets all the elements of the matrix *m* to the value *x*.

Function: void **gsl_matrix_set_zero** (*gsl_matrix * m*)

This function sets all the elements of the matrix *m* to zero.

Function: void **gsl_matrix_set_identity** (*gsl_matrix * m*)

This function sets the elements of the matrix *m* to the corresponding elements of the identity matrix, $m(i,j) = \delta(i,j)$, i.e. a unit diagonal with all off-diagonal elements zero. This applies to both square and rectangular matrices.

Reading and writing matrices

The library provides functions for reading and writing matrices to a file as binary data or formatted text.

Function: int **gsl_matrix_fwrite** (*FILE * stream, const gsl_matrix * m*)

This function writes the elements of the matrix *m* to the stream *stream* in binary format. The return value is 0 for success and `GSL_EFAILED` if there was a problem writing to the file. Since the data is written in

the native binary format it may not be portable between different architectures.

Function: `int gsl_matrix_fread (FILE * stream, gsl_matrix * m)`

This function reads into the matrix *m* from the open stream *stream* in binary format. The matrix *m* must be preallocated with the correct dimensions since the function uses the size of *m* to determine how many bytes to read. The return value is 0 for success and `GSL_EFAILED` if there was a problem reading from the file. The data is assumed to have been written in the native binary format on the same architecture.

Function: `int gsl_matrix_fprintf (FILE * stream, const gsl_matrix * m, const char * format)`

This function writes the elements of the matrix *m* line-by-line to the stream *stream* using the format specifier *format*, which should be one of the `%g`, `%e` or `%f` formats for floating point numbers and `%d` for integers. The function returns 0 for success and `GSL_EFAILED` if there was a problem writing to the file.

Function: `int gsl_matrix_fscanf (FILE * stream, gsl_matrix * m)`

This function reads formatted data from the stream *stream* into the matrix *m*. The matrix *m* must be preallocated with the correct dimensions since the function uses the size of *m* to determine how many numbers to read. The function returns 0 for success and `GSL_EFAILED` if there was a problem reading from the file.

Matrix views

A matrix view is a temporary object, stored on the stack, which can be used to operate on a subset of matrix elements. Matrix views can be defined for both constant and non-constant matrices using separate types that preserve constness. A matrix view has the type `gsl_matrix_view` and a constant matrix view has the type `gsl_matrix_const_view`. In both cases the elements of the view can be accessed using the `matrix` component of the view object. A pointer `gsl_matrix *` or `const gsl_matrix *` can be obtained by taking the address of the `matrix` component with the `&` operator. In addition to matrix views it is also possible to create vector views of a matrix, such as row or column views.

Function: `gsl_matrix_view gsl_matrix_submatrix (gsl_matrix * m, size_t k1, size_t k2, size_t n1, size_t n2)`

Function: `gsl_matrix_const_view gsl_matrix_const_submatrix (const gsl_matrix * m, size_t k1, size_t k2, size_t n1, size_t n2)`

These functions return a matrix view of a submatrix of the matrix *m*. The upper-left element of the submatrix is the element (*k1*,*k2*) of the original matrix. The submatrix has *n1* rows and *n2* columns. The physical number of columns in memory given by *tda* is unchanged. Mathematically, the (*i*,*j*)-th element of the new matrix is given by,

$$m'(i, j) = m \rightarrow \text{data}[(k1 * m \rightarrow \text{tda} + k1) + i * m \rightarrow \text{tda} + j]$$

where the index *i* runs from 0 to *n1*-1 and the index *j* runs from 0 to *n2*-1.

The `data` pointer of the returned matrix struct is set to null if the combined parameters (*i*,*j*,*n1*,*n2*,*tda*) overrun the ends of the original matrix.

The new matrix view is only a view of the block underlying the existing matrix, *m*. The block containing the elements of *m* is not owned by the new matrix view. When the view goes out of scope the original matrix *m* and its block will continue to exist. The original memory can only be deallocated by freeing the original matrix. Of course, the original matrix should not be deallocated while the view is still in use.

The function `gsl_matrix_const_submatrix` is equivalent to `gsl_matrix_submatrix` but can be used for matrices which are declared `const`.

Function: `gsl_matrix_view` **`gsl_matrix_view_array`** (*double * base, size_t n1, size_t n2*)

Function: `gsl_matrix_const_view` **`gsl_matrix_const_view_array`** (*const double * base, size_t n1, size_t n2*)

These functions return a matrix view of the array *base*. The matrix has *n1* rows and *n2* columns. The physical number of columns in memory is also given by *n2*. Mathematically, the (i,j) -th element of the new matrix is given by,

$$m'(i,j) = \text{base}[i*n2 + j]$$

where the index *i* runs from 0 to *n1*-1 and the index *j* runs from 0 to *n2*-1.

The new matrix is only a view of the array *base*. When the view goes out of scope the original array *base* will continue to exist. The original memory can only be deallocated by freeing the original array. Of course, the original array should not be deallocated while the view is still in use.

The function `gsl_matrix_const_view_array` is equivalent to `gsl_matrix_view_array` but can be used for matrices which are declared `const`.

Function: `gsl_matrix_view` **`gsl_matrix_view_array_with_tda`** (*double * base, size_t n1, size_t n2, size_t tda*)

Function: `gsl_matrix_const_view` **`gsl_matrix_const_view_array_with_tda`** (*const double * base, size_t n1, size_t n2, size_t tda*)

These functions return a matrix view of the array *base* with a physical number of columns *tda* which may differ from corresponding the dimension of the matrix. The matrix has *n1* rows and *n2* columns, and the physical number of columns in memory is given by *tda*. Mathematically, the (i,j) -th element of the new matrix is given by,

$$m'(i,j) = \text{base}[i*tda + j]$$

where the index *i* runs from 0 to *n1*-1 and the index *j* runs from 0 to *n2*-1.

The new matrix is only a view of the array *base*. When the view goes out of scope the original array *base* will continue to exist. The original memory can only be deallocated by freeing the original array. Of course, the original array should not be deallocated while the view is still in use.

The function `gsl_matrix_const_view_array_with_tda` is equivalent to `gsl_matrix_view_array_with_tda` but can be used for matrices which are declared `const`.

Function: `gsl_matrix_view` **`gsl_matrix_view_vector`** (*gsl_vector * v, size_t n1, size_t n2*)

Function: `gsl_matrix_const_view` **`gsl_matrix_const_view_vector`** (*const gsl_vector * v, size_t n1, size_t n2*)

These functions return a matrix view of the vector *v*. The matrix has *n1* rows and *n2* columns. The vector must have unit stride. The physical number of columns in memory is also given by *n2*. Mathematically, the (i,j) -th element of the new matrix is given by,

$$m'(i,j) = v->\text{data}[i*n2 + j]$$

where the index *i* runs from 0 to *n1*-1 and the index *j* runs from 0 to *n2*-1.

The new matrix is only a view of the vector v . When the view goes out of scope the original vector v will continue to exist. The original memory can only be deallocated by freeing the original vector. Of course, the original vector should not be deallocated while the view is still in use.

The function `gsl_matrix_const_view_vector` is equivalent to `gsl_matrix_view_vector` but can be used for matrices which are declared `const`.

Function: `gsl_matrix_view` **`gsl_matrix_view_vector_with_tda`** (*`gsl_vector * v, size_t n1, size_t n2, size_t tda`*)
Function: `gsl_matrix_const_view` **`gsl_matrix_const_view_vector_with_tda`** (*`const gsl_vector * v, size_t n1, size_t n2, size_t tda`*)

These functions return a matrix view of the vector v with a physical number of columns tda which may differ from the corresponding matrix dimension. The vector must have unit stride. The matrix has $n1$ rows and $n2$ columns, and the physical number of columns in memory is given by tda . Mathematically, the (i,j) -th element of the new matrix is given by,

$$m'(i,j) = v->data[i*tda + j]$$

where the index i runs from 0 to $n1-1$ and the index j runs from 0 to $n2-1$.

The new matrix is only a view of the vector v . When the view goes out of scope the original vector v will continue to exist. The original memory can only be deallocated by freeing the original vector. Of course, the original vector should not be deallocated while the view is still in use.

The function `gsl_matrix_const_view_vector_with_tda` is equivalent to `gsl_matrix_view_vector_with_tda` but can be used for matrices which are declared `const`.

Creating row and column views

In general there are two ways to access an object, by reference or by copying. The functions described in this section create vector views which allow access to a row or column of a matrix by reference. Modifying elements of the view is equivalent to modifying the matrix, since both the vector view and the matrix point to the same memory block.

Function: `gsl_vector_view` **`gsl_matrix_row`** (*`gsl_matrix * m, size_t i`*)

Function: `gsl_vector_const_view` **`gsl_matrix_const_row`** (*`const gsl_matrix * m, size_t i`*)

These functions return a vector view of the i -th row of the matrix m . The `data` pointer of the new vector is set to null if i is out of range.

The function `gsl_vector_const_row` is equivalent to `gsl_matrix_row` but can be used for matrices which are declared `const`.

Function: `gsl_vector_view` **`gsl_matrix_column`** (*`gsl_matrix * m, size_t j`*)

Function: `gsl_vector_const_view` **`gsl_matrix_const_column`** (*`const gsl_matrix * m, size_t j`*)

These functions return a vector view of the j -th column of the matrix m . The `data` pointer of the new vector is set to null if j is out of range.

The function `gsl_vector_const_column` equivalent to `gsl_matrix_column` but can be used for matrices which are declared `const`.

Function: `gsl_vector_view` **`gsl_matrix_diagonal`** (*`gsl_matrix * m`*)

Function: `gsl_vector_const_view` **`gsl_matrix_const_diagonal`** (*`const gsl_matrix * m`*)

These functions return a vector view of the diagonal of the matrix *m*. The matrix *m* is not required to be square. For a rectangular matrix the length of the diagonal is the same as the smaller dimension of the matrix.

The function `gsl_matrix_const_diagonal` is equivalent to `gsl_matrix_diagonal` but can be used for matrices which are declared `const`.

Function: `gsl_vector_view` **`gsl_matrix_subdiagonal`** (*`gsl_matrix * m, size_t k`*)

Function: `gsl_vector_const_view` **`gsl_matrix_const_subdiagonal`** (*`const gsl_matrix * m, size_t k`*)

These functions return a vector view of the *k*-th subdiagonal of the matrix *m*. The matrix *m* is not required to be square. The diagonal of the matrix corresponds to *k* = 0.

The function `gsl_matrix_const_subdiagonal` is equivalent to `gsl_matrix_subdiagonal` but can be used for matrices which are declared `const`.

Function: `gsl_vector_view` **`gsl_matrix_superdiagonal`** (*`gsl_matrix * m, size_t k`*)

Function: `gsl_vector_const_view` **`gsl_matrix_const_superdiagonal`** (*`const gsl_matrix * m, size_t k`*)

These functions return a vector view of the *k*-th superdiagonal of the matrix *m*. The matrix *m* is not required to be square. The diagonal of the matrix corresponds to *k* = 0.

The function `gsl_matrix_const_superdiagonal` is equivalent to `gsl_matrix_superdiagonal` but can be used for matrices which are declared `const`.

Copying matrices

Function: `int` **`gsl_matrix_memcpy`** (*`gsl_matrix * dest, const gsl_matrix * src`*)

This function copies the elements of the matrix *src* into the matrix *dest*. The two matrices must have the same size.

Function: `int` **`gsl_matrix_swap`** (*`gsl_matrix * m1, const gsl_matrix * m2`*)

This function exchanges the elements of the matrices *m1* and *m2* by copying. The two matrices must have the same size.

Copying rows and columns

The functions described in this section copy a row or column of a matrix into a vector. This allows the elements of the vector and the matrix to be modified independently. Note that if the matrix and the vector point to overlapping regions of memory then the result will be undefined. The same effect can be achieved with more generality using `gsl_vector_memcpy` with vector views of rows and columns.

Function: `int` **`gsl_matrix_get_row`** (*`gsl_vector * v, const gsl_matrix * m, size_t i`*)

This function copies the elements of the *i*-th row of the matrix *m* into the vector *v*. The length of the vector must be the same as the length of the row.

Function: `int` **`gsl_matrix_get_col`** (*`gsl_vector * v, const gsl_matrix * m, size_t j`*)

This function copies the elements of the i -th column of the matrix m into the vector v . The length of the vector must be the same as the length of the column.

Function: `int gsl_matrix_set_row (gsl_matrix * m, size_t i, const gsl_vector * v)`

This function copies the elements of the vector v into the i -th row of the matrix m . The length of the vector must be the same as the length of the row.

Function: `int gsl_matrix_set_col (gsl_matrix * m, size_t j, const gsl_vector * v)`

This function copies the elements of the vector v into the j -th column of the matrix m . The length of the vector must be the same as the length of the column.

Exchanging rows and columns

The following functions can be used to exchange the rows and columns of a matrix.

Function: `int gsl_matrix_swap_rows (gsl_matrix * m, size_t i, size_t j)`

This function exchanges the i -th and j -th rows of the matrix m in-place.

Function: `int gsl_matrix_swap_columns (gsl_matrix * m, size_t i, size_t j)`

This function exchanges the i -th and j -th columns of the matrix m in-place.

Function: `int gsl_matrix_swap_rowcol (gsl_matrix * m, size_t i, size_t j)`

This function exchanges the i -th row and j -th column of the matrix m in-place. The matrix must be square for this operation to be possible.

Function: `int gsl_matrix_transpose_memcpy (gsl_matrix * dest, gsl_matrix * src)`

This function makes the matrix $dest$ the transpose of the matrix src by copying the elements of src into $dest$. This function works for all matrices provided that the dimensions of the matrix $dest$ match the transposed dimensions of the matrix src .

Function: `int gsl_matrix_transpose (gsl_matrix * m)`

This function replaces the matrix m by its transpose by copying the elements of the matrix in-place. The matrix must be square for this operation to be possible.

Matrix operations

The following operations are only defined for real matrices.

Function: `int gsl_matrix_add (gsl_matrix * a, const gsl_matrix * b)`

This function adds the elements of matrix b to the elements of matrix a , $a'(i,j) = a(i,j) + b(i,j)$. The two matrices must have the same dimensions.

Function: `int gsl_matrix_sub (gsl_matrix * a, const gsl_matrix * b)`

This function subtracts the elements of matrix b from the elements of matrix a , $a'(i,j) = a(i,j) - b(i,j)$. The two matrices must have the same dimensions.

Function: `int gsl_matrix_mul_elements (gsl_matrix * a, const gsl_matrix * b)`

This function multiplies the elements of matrix a by the elements of matrix b , $a'(i,j) = a(i,j) * b(i,j)$. The two matrices must have the same dimensions.

Function: `int gsl_matrix_div_elements (gsl_matrix * a, const gsl_matrix * b)`

This function divides the elements of matrix *a* by the elements of matrix *b*, $a'(i,j) = a(i,j) / b(i,j)$. The two matrices must have the same dimensions.

Function: `int gsl_matrix_scale (gsl_matrix * a, const double x)`

This function multiplies the elements of matrix *a* by the constant factor *x*, $a'(i,j) = x a(i,j)$.

Function: `int gsl_matrix_add_constant (gsl_matrix * a, const double x)`

This function adds the constant value *x* to the elements of the matrix *a*, $a'(i,j) = a(i,j) + x$.

Finding maximum and minimum elements of matrices

Function: `double gsl_matrix_max (const gsl_matrix * m)`

This function returns the maximum value in the matrix *m*.

Function: `double gsl_matrix_min (const gsl_matrix * m)`

This function returns the minimum value in the matrix *m*.

Function: `void gsl_matrix_minmax (const gsl_matrix * m, double * min_out, double * max_out)`

This function returns the minimum and maximum values in the matrix *m*, storing them in *min_out* and *max_out*.

Function: `void gsl_matrix_max_index (const gsl_matrix * m, size_t * imax, size_t * jmax)`

This function returns the indices of the maximum value in the matrix *m*, storing them in *imax* and *jmax*. When there are several equal maximum elements then the first element found is returned.

Function: `void gsl_matrix_min_index (const gsl_matrix * m, size_t * imax, size_t * jmax)`

This function returns the indices of the minimum value in the matrix *m*, storing them in *imax* and *jmax*. When there are several equal minimum elements then the first element found is returned.

Function: `void gsl_matrix_minmax_index (const gsl_matrix * m, size_t * imin, size_t * imax)`

This function returns the indices of the minimum and maximum values in the matrix *m*, storing them in (*imin*, *jmin*) and (*imax*, *jmax*). When there are several equal minimum or maximum elements then the first elements found are returned.

Matrix properties

Function: `int gsl_matrix_isnull (const gsl_matrix * m)`

This function returns 1 if all the elements of the matrix *m* are zero, and 0 otherwise.

Example programs for matrices

The program below shows how to allocate, initialize and read from a matrix using the functions `gsl_matrix_alloc`, `gsl_matrix_set` and `gsl_matrix_get`.

```
#include <stdio.h>
#include <gsl/gsl_matrix.h>
```

```

int
main (void)
{
    int i, j;
    gsl_matrix * m = gsl_matrix_alloc (10, 3);

    for (i = 0; i < 10; i++)
        for (j = 0; j < 3; j++)
            gsl_matrix_set (m, i, j, 0.23 + 100*i + j);

    for (i = 0; i < 100; i++)
        for (j = 0; j < 3; j++)
            printf("m(%d,%d) = %g\n", i, j,
                gsl_matrix_get (m, i, j));

    return 0;
}

```

Here is the output from the program. The final loop attempts to read outside the range of the matrix `m`, and the error is trapped by the range-checking code in `gsl_matrix_get`.

```

m(0,0) = 0.23
m(0,1) = 1.23
m(0,2) = 2.23
m(1,0) = 100.23
m(1,1) = 101.23
m(1,2) = 102.23
...
m(9,2) = 902.23
gsl: matrix_source.c:13: ERROR: first index out of range
IOT trap/Abort (core dumped)

```

The next program shows how to write a matrix to a file.

```

#include <stdio.h>
#include <gsl/gsl_matrix.h>

int
main (void)
{
    int i, j, k = 0;
    gsl_matrix * m = gsl_matrix_alloc (100, 100);
    gsl_matrix * a = gsl_matrix_alloc (100, 100);

    for (i = 0; i < 100; i++)
        for (j = 0; j < 100; j++)
            gsl_matrix_set (m, i, j, 0.23 + i + j);

    {
        FILE * f = fopen("test.dat", "w");
        gsl_matrix_fwrite (f, m);
        fclose (f);
    }

    {
        FILE * f = fopen("test.dat", "r");

```

```

    gsl_matrix_fread (f, a);
    fclose (f);
}

for (i = 0; i < 100; i++)
    for (j = 0; j < 100; j++)
    {
        double mij = gsl_matrix_get(m, i, j);
        double aij = gsl_matrix_get(a, i, j);
        if (mij != aij) k++;
    }

printf("differences = %d (should be zero)\n", k);
return (k > 0);
}

```

After running this program the file `test.dat` should contain the elements of m , written in binary format. The matrix which is read back in using the function `gsl_matrix_fread` should be exactly equal to the original matrix.

The following program demonstrates the use of vector views. The program computes the column-norms of a matrix.

```

#include <math.h>
#include <stdio.h>
#include <gsl/gsl_matrix.h>
#include <gsl/gsl_blas.h>

int
main (void)
{
    size_t i,j;

    gsl_matrix *m = gsl_matrix_alloc (10, 10);

    for (i = 0; i < 10; i++)
        for (j = 0; j < 10; j++)
            gsl_matrix_set (m, i, j, sin (i) + cos (j));

    for (j = 0; j < 10; j++)
    {
        gsl_vector_view column = gsl_matrix_column (m, j);
        double d;

        d = gsl_blas_dnrm2 (&column.vector);

        printf ("matrix column %d, norm = %g\n", j, d);
    }

    gsl_matrix_free (m);
}

```

Here is the output of the program, which can be confirmed using GNU OCTAVE,

```

$ ./a.out
matrix column 0, norm = 4.31461

```

```
matrix column 1, norm = 3.1205
matrix column 2, norm = 2.19316
matrix column 3, norm = 3.26114
matrix column 4, norm = 2.53416
matrix column 5, norm = 2.57281
matrix column 6, norm = 4.20469
matrix column 7, norm = 3.65202
matrix column 8, norm = 2.08524
matrix column 9, norm = 3.07313

octave> m = sin(0:9)' * ones(1,10)
          + ones(10,1) * cos(0:9);
octave> sqrt(sum(m.^2))
ans =

    4.3146    3.1205    2.1932    3.2611    2.5342    2.5728
    4.2047    3.6520    2.0852    3.0731
```

References and Further Reading

The block, vector and matrix objects in GSL follow the `valarray` model of C++. A description of this model can be found in the following reference,

- B. Stroustrup, *The C++ Programming Language* (3rd Ed), Section 22.4 Vector Arithmetic. Addison-Wesley 1997, ISBN 0-201-88954-4.

Go to the [first](#), [previous](#), [next](#), [last](#) section, [table of contents](#).